

AI Engineer

Engineering — AI/ML Platform Team | Full-time | Hybrid

About the Role

We're looking for an AI Engineer to build, train, and deploy machine learning and generative AI systems in production. You'll work closely with product and platform teams to turn research and prototypes into reliable, real-world services.

What You'll Do

- Design, train, and evaluate ML/deep learning models for production use cases
- Build data pipelines and feature engineering workflows
- Deploy models as scalable services and monitor them in production
- Work with LLMs: prompt engineering, RAG, fine-tuning, and agentic workflows
- Collaborate with product and engineering to turn requirements into ML solutions

What We're Looking For

- Bachelor's degree in CS, Data Science, or related field (or equivalent experience)
- Proficiency in Python and at least one ML framework (PyTorch or TensorFlow)
- Solid grounding in data structures, algorithms, and computer architecture
- Experience with LLM frameworks (LangGraph/LangChain), RAG, or vector databases is a plus
- Familiarity with cloud ML infrastructure and MLOps practices is a plus